

Guangzhi (Allen) Su

+1 (919)-641-9352 | allens@andrew.cmu.edu | <https://guangzhisu.github.io/>

EDUCATION

Carnegie Mellon University

MS in Artificial Intelligence and Innovation

Pittsburgh, PA

Expected: May 2028

Duke University

BS in Applied Mathematics & Computer Science | GPA: 3.85/4.0 | Magna Cum Laude

Durham, NC

May 2026

SKILLS

Programming & Scripting

Python, Java, C++, TypeScript, JavaScript, SQL, HTML, CSS, R

ML & AI

PyTorch, Numpy, Pandas, Huggingface, DeepSpeed, LangChain, LangGraph, Docker

System & Web

Docker, Linux, Git, Next.js, React, Electron, Ollama

Language

English (fluent), Mandarin (native), Spanish (entry)

PROFESSIONAL EXPERIENCE

AI Product Manager Intern

May 2026 – Aug 2026

TikTok (Bytedance)

Beijing, China

- Co-designed a two-stage agentic moderation pipeline in which a routing agent selects the top-5 candidate categories before policy-specific decision-tree agents sequentially evaluate natural-language rules and violation signals to human reviewers; optimized ~10 production prompts through automated root cause analysis and manual iteration over 10,000+ cases, achieved 32% increase on recall, 3.7% increase on precision, and 22% on F1 score during offline experiment.
- Built evaluation frameworks for JIMU AIMate, a low-code webpage template development platform with ~1,000-users, benchmarking AI-generated task templates against ~15 PRD-ground-truth pairs and categorizing failures across different reasons; achieved an 85% generation success rate and 50% usage rate, with non-engineers surpassing engineers in template creation volume for the first time.

Founder & Project Lead

Dec 2025 – Present

Prospectus.AI ([Website](#))

Suzhou, China

- Developed and deployed an agentic AI platform with two core systems: a prospectus drafting engine that reverse-engineers existing prospectus into section-level rules for RAG-based generation from multimodal company documents, and an IPO diagnostic engine combining AI-based information extraction with deterministic thresholds and qualitative assessment across multiple markets.
- Established ground-truth datasets, prompt/data inspection tools, and automated evaluation pipelines for generation error attribution and iterative refinement; deployed the product on Vercel, secured ¥25,000 in funding support, demoed it to three companies, and won 1st Place in the Finance Track at HackDKU 2026.

Technology Research and Development Intern

Dec 2023 – Jan 2024

Microsoft-INESA AI Innovation Center

Shanghai, China

- Produced technical landscape reports on frontier AI research and industry developments, synthesizing academic and industry sources on AI-generated content detection, LLM training and fine-tuning, and emerging model capabilities across China and global markets; compared representative approaches, technical trade-offs, and application scenarios.
- Evaluated end-to-end LLM development and deployment workflows, covering model fine-tuning, Retrieval-Augmented Generation (RAG), Docker-based containerization, and cloud-based inference infrastructure; analyzed how different system configurations support model adaptation, deployment, and scalable serving.

RESEARCH EXPERIENCE

Research Assistant - Multimodal Clinical Risk Prediction System

Jan 2026 – Present

Duke University School of Medicine, Advisor: Prof. Rishi Kamaleswaran, Ph.D.

Durham, NC, USA

- Developed a multimodal clinical forecasting framework that predicts a patient's future chest X-ray representation from a prior CXR and intervening ECG waveforms, providing an intermediate representation of future clinical state for downstream risk prediction.
- Designed and trained a JEPA-style multimodal model with cross-attention between CXR and ECG representations. Reduced future-CXR latent MSE by 9.3% and achieved 0.794 Macro-AUROC and 0.866 Macro-AUPRC on six-disease future-status classification.

Research Assistant - Agent-Based Framework for Efficient Long-Video Understanding ([Github](#))

Jun 2024 – Jul 2025

Duke University, Advisor: Prof. Yiran Chen, Ph.D.

Durham, NC, USA

- Proposed and implemented a parallel multi-agent system for long-video understanding, integrating MoE architecture and specialized agents to process extended video sequences with 40% cut for inference time compared to SOTA method.
- Built the system using Huggingface pipeline with multi-processing for asynchronous execution, achieving 60% accuracy on several first-person benchmarks, demonstrating competitive performance and substantial speedups.

PUBLICATIONS

- G. Su et al. SmoothGuard: Defending Multimodal Large Language Models (MLLMs) with Noise Perturbation and Clustering Aggregation.** IEEE ICDM 2025 ([Paper](#))
- Y. Liu et al. Keyframe-oriented vision token pruning: Enhancing efficiency of large vision language models on long-form video processing.** ICCV 2025 ([Paper](#))